

Layer sweep: OLS activation → function-vector regression, full-dim test MSE vs layer (activation PCs=32, FV PCs=16; MSE in 4096-d FV space; fit on train tasks → eval on held-out test tasks) (layers <1 omitted: embedding layer MSE is off-scale)

