

Layer sweep: OLS activation → function-vector regression, test MSE vs layer (k=5, fit on train tasks → eval on held-out test tasks) (layers <1 omitted: embedding layer MSE is off-scale)

